

3. Short Writeup

3.1 Estimate of two coins are: - $\theta_A = 0.568$ and $\theta_B = 0.494$

3.2 There were 4 ways to make multiple API calls to the given URL. 1- Use `requests.get (URL)` in a for loop, 2- Use `asyncio` to make asynchronous call to the given URL, 3- Use `threads`, 4- Use `multiprocessing`.

The first option is clearly a brute force approach with high runtime. I chose option 2 over 3 and 4 because the task in hand deals with doing many connections. Moreover, due to GIL usage of threads in python does not achieve true parallelism. The task is high I/O bound and hence using `multiprocessing` is not feasible.

After running `asyncio` to make 30 API calls I extract the "data" from the response json. These draws are then appended to a list. Since the coin tosses are un-labelled, this is a semi-supervised learning problem. I assumed the data is Gaussian Fit data. I used `sklearn.mixture.gmm` to fit the data as follows.

```
gmm = mixture.GaussianMixture (n_components=2, max_iter=10000, covariance_type='tied',  
init_params = "random").fit (coin_flips)
```

I had to choose the covariance type variable in the model, I chose "tied" instead of "full" because "full" gives good results if we have a large dataset, in smaller datasets like the one in hand tends to overfit the data.

I had to choose "random" for `init_params`, we don't know that our data is clustered via `kmeans` algorithm or not, assuming it is random I chose "random" for `init_params`. If we know the coin draws are `kmeans` clustered then choosing "kmeans" would be appropriate.

This gives us the model with the estimated gaussian parameters - mean, variances etc. The labels are then predicted for each coin draw using the fitted model. I get a dictionary with 0 and 1 as the keys which are the type of each coin, and coin draws as values. Now that I have labelled data, theta values are calculated.